

# CSCE 421 Final Project: AUC Maximization by Deep Learning for Imbalanced Medical Image Classification

Jonathan Zhao, David-Tyler Ighedosa, Kyumin Lee

## Abstract

During the course of this class we've used different techniques and methods in the hopes of better optimizing machine learning models. In this project we combine these techniques in order to maximize the Area under the Curve which measures the performance of the machine learning model. These techniques were used on the Imbalanced Medical Image Classification that was provided.

## Introduction

Using ResNet18, LibAUC, MedMNIST, and PneumoniaMNIST different techniques were implemented in order to optimize our models. The three main methodologies we used were data augmentation, adjustment of the learning rate, controlling overfitting, and changing the optimizer. With these techniques we trained a deep learning model to maximize the area under the curve for imbalanced medical images.

## Methods

In this study, we explore several methods for enhancing the performance of machine learning models:

### Data Augmentation

We performed many data augmentation strategies before including L1 and L2 regularization. Data has been provided with no augmentation, some augmentation, and the final augmentation we tried. The first 10 epochs and their respective AUC values were provided that demonstrate the effectiveness of the data augmentation technique used. From the following data we can see that the AUC value does not increase with any change, some of the changes we made provided

lower AUC values. Changes that were performed include flipping, random rotation, color jitter, and blur.

### **Adjusting the Learning Rate**

The learning rate affects how fast the model converges during training. Having a learning rate that's too low will slow down the convergence, while having a learning rate that's too high will cause the model to overshoot, and become unstable. We adjusted the learning rate from the default of 0.1 to 0.2, which helped speed up the convergence at the lower epochs, while also not being too high that it will cause the model to overshoot.

### **Controlling Overfitting**

- L1 Regularization
  - L1 regularization adds a penalty proportional to the absolute value of the coefficients to the loss function. It encourages sparsity in the model, meaning it tends to push the weights of less important features to exactly zero. L1 Regularization can simplify the model and make it more interpretable by selecting only the most important features. It can also help prevent overfitting by reducing the model's complexity.
- L2 Regularization
  - L2 Regularization, also known as weight decay, adds a penalty proportional to the square of the coefficients to the loss function. It tends to shrink the weights of the model overall without necessarily setting them to zero. L2 Regularization is used to prevent overfitting by discouraging overly complex models. It smooths the model's learned parameters and can improve generalization to unseen data.

### **Changing Optimizers**

We tried changing the optimizer to using the Adam optimizer. The Adam optimizer is an adaptive learning rate optimization algorithm that combines the advantages of two other popular optimization algorithms, AdaGrad and RMSProp. Going into it we did not know how much it would affect the results but from our experiments we could see that it had a decent impact compared to other optimization methods.

## Experiments and Data

### No Data Augmentation without Optimizers:

Epoch: 0; Test AUC: 0.38784461152882205  
Epoch: 1; Test AUC: 0.5352965747702589  
Epoch: 2; Test AUC: 0.7717209690893903  
Epoch: 3; Test AUC: 0.47556390977443613  
Epoch: 4; Test AUC: 0.7017543859649124  
Epoch: 5; Test AUC: 0.5639097744360902  
Epoch: 6; Test AUC: 0.7758980785296575  
Epoch: 7; Test AUC: 0.7926065162907268  
Epoch: 8; Test AUC: 0.7731829573934836  
Epoch: 9; Test AUC: 0.7909356725146199

```
train_transform_original = transforms.Compose([
    transforms.ToPILImage(),
    transforms.Grayscale(3),
    transforms.Resize((32, 32)),
    transforms.ToTensor(),
    transforms.Normalize(0.5, 0.5),
])
```

### Some Data Augmentation without Optimizers:

Epoch: 0; Test AUC: 0.3504594820384294  
Epoch: 1; Test AUC: 0.39598997493734334  
Epoch: 2; Test AUC: 0.3834586466165414  
Epoch: 3; Test AUC: 0.38345864661654133  
Epoch: 4; Test AUC: 0.31370091896407687  
Epoch: 5; Test AUC: 0.37259816207184626  
Epoch: 6; Test AUC: 0.5348788638262322  
Epoch: 7; Test AUC: 0.7224310776942356  
Epoch: 8; Test AUC: 0.7428989139515456  
Epoch: 9; Test AUC: 0.6917293233082707

```
train_transform_temp = transforms.Compose([
    transforms.ToPILImage(),
    transforms.Grayscale(3),
    transforms.RandomHorizontalFlip(p=1), #
    transforms.RandomRotation(degrees=10), #
    transforms.ColorJitter(brightness=40),
    transforms.Resize((32, 32)),
    transforms.ToTensor(),
    transforms.Normalize(0.5, 0.5)
])
```

### Final data Augmentation without Optimizers:

Epoch: 0; Test AUC: 0.3767752715121136  
Epoch: 1; Test AUC: 0.5421888053467001  
Epoch: 2; Test AUC: 0.5954469507101086  
Epoch: 3; Test AUC: 0.6351294903926483  
Epoch: 4; Test AUC: 0.7007101086048455  
Epoch: 5; Test AUC: 0.72890559732665  
Epoch: 6; Test AUC: 0.6313700918964077  
Epoch: 7; Test AUC: 0.6349206349206349  
Epoch: 8; Test AUC: 0.6449456975772766  
Epoch: 9; Test AUC: 0.3845029239766082

```
train_transform_final = transforms.Compose([
    transforms.ToPILImage(),
    transforms.Grayscale(3),
    transforms.RandomHorizontalFlip(1), # Random horizon
    transforms.RandomVerticalFlip(1),
    transforms.RandomRotation(degrees=(60,61)), #random
    transforms.ColorJitter(brightness=40, contrast=45,
                           saturation=100, hue=0.3),
    transforms.GaussianBlur(5, sigma=(0.2, 3.0)),
    transforms.Resize((32, 32)), |
    transforms.ToTensor(),
    transforms.Normalize(0.5, 0.5)
])
```

These Optimization Tests were done with the following data augmentations:

### No Optimizations:

Epoch: 0; Test AUC: 0.34398496240601506  
Epoch: 1; Test AUC: 0.42251461988304095  
Epoch: 2; Test AUC: 0.6027568922305765  
Epoch: 3; Test AUC: 0.47368421052631576  
Epoch: 4; Test AUC: 0.730576441102757  
Epoch: 5; Test AUC: 0.3742690058479532  
Epoch: 6; Test AUC: 0.4880952380952381  
Epoch: 7; Test AUC: 0.6706349206349206  
Epoch: 8; Test AUC: 0.6409774436090225  
Epoch: 9; Test AUC: 0.6409774436090226

```
train_transform = transforms.Compose([
    transforms.ToPILImage(),
    transforms.Grayscale(3),
    transforms.RandomRotation(degrees=10),
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomVerticalFlip(p=0.5),
    transforms.Resize((32, 32)),
    transforms.ToTensor(),
    transforms.Normalize(0.5, 0.5),
])
```


### Controlling Overfitting with L1 Regularization

#### Using K fold Cross Validation:

Optimizing with L1 Regularization

Epoch: 0; Test AUC: 0.6216791979949875  
Epoch: 1; Test AUC: 0.6458020050125314  
Epoch: 2; Test AUC: 0.5020050125313282  
Epoch: 3; Test AUC: 0.6760651629072683  
Epoch: 4; Test AUC: 0.6307226399331662  
Epoch: 5; Test AUC: 0.6573934837092732  
Epoch: 6; Test AUC: 0.643107769423559  
Epoch: 7; Test AUC: 0.5489348370927318  
Epoch: 8; Test AUC: 0.5925020885547202  
Epoch: 9; Test AUC: 0.6340434419381789

#### No Optimization and L1 Regularization

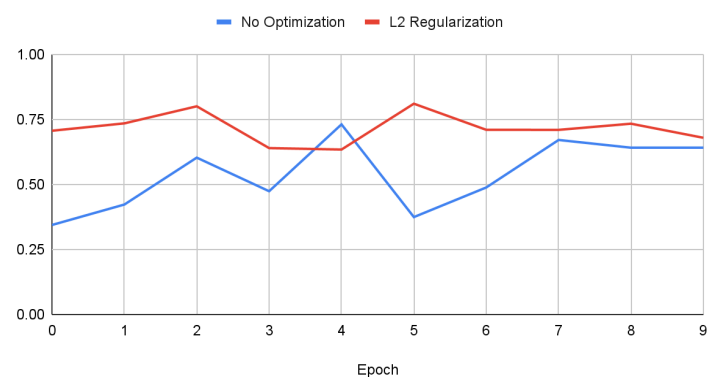


### Controlling Overfitting with L2 Regularization

#### Using K fold Cross Validation:

Epoch: 0; Test AUC: 0.7060985797827902  
Epoch: 1; Test AUC: 0.7344611528822055  
Epoch: 2; Test AUC: 0.8000417710944028  
Epoch: 3; Test AUC: 0.639390142021721  
Epoch: 4; Test AUC: 0.6338345864661654  
Epoch: 5; Test AUC: 0.8100250626566415  
Epoch: 6; Test AUC: 0.7096700083542189  
Epoch: 7; Test AUC: 0.7094820384294069  
Epoch: 8; Test AUC: 0.7331662489557227  
Epoch: 9; Test AUC: 0.679030910609858


#### No Optimization and L2 Regularization



### Using the Adam Optimizer with 0.2 learning rate:

Epoch: 0; Test AUC: 0.36737677527151213  
Epoch: 1; Test AUC: 0.4066416040100251  
Epoch: 2; Test AUC: 0.3391812865497076  
Epoch: 3; Test AUC: 0.4733709273182957  
Epoch: 4; Test AUC: 0.5110693400167083  
Epoch: 5; Test AUC: 0.5572263993316624  
Epoch: 6; Test AUC: 0.46825396825396826  
Epoch: 7; Test AUC: 0.43859649122807015  
Epoch: 8; Test AUC: 0.4951963241436925  
Epoch: 9; Test AUC: 0.424812030075188


No Optimization and Adam Optimizer with 0.2 Learning Rate



### Using the Adam Optimizer with 0.1 learning rate:

Epoch: 0; Test AUC: 0.38053467000835417  
Epoch: 1; Test AUC: 0.4502923976608187  
Epoch: 2; Test AUC: 0.3778195488721805  
Epoch: 3; Test AUC: 0.40517961570593153  
Epoch: 4; Test AUC: 0.526733500417711  
Epoch: 5; Test AUC: 0.5190058479532164  
Epoch: 6; Test AUC: 0.49122807017543857  
Epoch: 7; Test AUC: 0.5309106098579782  
Epoch: 8; Test AUC: 0.5104427736006684  
Epoch: 9; Test AUC: 0.5043859649122807

No Optimization and Adam Optimizer with 0.1 Learning Rate



## Conclusions

With L1 Regularization, we can see that the results can vary from epoch to epoch to some performing better with L1 Regularization compared to the original data. But on average, we do see a slight increase in performance due to the optimization. On the other hand with L2 Regularization we see a general increase in performance for almost every epoch showing an increase in performance due to the optimization on average. For changing the optimizer with the two different learning rates we see a general decrease in results as compared to the PESG optimizer that was originally used.

With the learning rate, the base model uses a value of 0.1, which provides adequate AUC values. Increasing the learning rate to 0.5 caused the AUC values of the smaller value epochs to be better, but the higher value epochs had overall worse AUC values. Overall, the best AUC values were obtained with a learning rate of 0.2.

With Data Augmentation, we can conclude that often some changes to the data might lower the performance of the model. Some of the epochs saw noticeable improvements with augmentation but as a whole some of the changes that were made to the data led to lower AUC performance. Further tests with varying levels of brightness, rotation, cropping, etc would provide a higher AUC.